

**CSS-587 Final Project Proposal:**  
**LP-SIFT Implementation and Optimization**  
David Li, Ben Schipunov, Kris Yu  
November 3, 2025

## 1. Selected Paper

**Title:** Local-peak scale-invariant feature transform for fast and random image stitching

**Authors:** Hao Li, Lipo Wang, Tianyun Zhao, Wei Zhao

**ArXiv:** <https://arxiv.org/pdf/2405.08578>

**MATLAB source:**

<https://github.com/haohaohao111222/LP-SIFT-Algorithm-Combined-with-RANSAC>

## 2. Project Description

### **Background:**

LP-SIFT replaces SIFT's computationally expensive Gaussian pyramid with multi-scale local peak detection, achieving 10-100× speedup on large images while maintaining stitching quality. The paper demonstrates successful stitching of nine 2600×1600 pixel images in under 160 seconds.

### **Proposal Summary:**

Validate paper results by replicating work using the provided dataset and MATLAB implementation on our hardware. Discuss and report results. Port work to C++, this will include reimplementing the paper's algorithm in C++ using OpenCV and optimizations on the Local Peak (LP) implementation. We'll rerun the benchmarking with the new C++ implementations and compare results, noting performance improvements. If time permits, our stretch goal would be further optimization leveraging GPU acceleration.

In conclusion for our project, we will provide a demo application that uses LP-SIFT to show first hand how fast the image stitching process improves, specifically performing best with large resolutions.

### **Goals**

1. Validate paper results using provided MATLAB implementation
2. Port LP-SIFT to C++ with OpenCV (library + CLI tool)
3. Benchmark against SIFT, ORB, BRISK, SURF on small/medium/large images
4. Apply CPU optimizations for local peak algorithm
5. Implement multi-image mosaic stitching for random fragments (demo application)
6. **Stretch:** GPU acceleration with CUDA

### 3. Input and Output

**Input:** Image pairs/sets in JPEG/PNG format across three size categories: small (~600×400), medium (1080×1920), and large (3072×4096). Parameters include interrogation window sizes L and matching threshold  $\Delta s$ .

**Output:** Stitched panoramic images, performance metrics (timing, feature counts, match quality), comparison visualizations, and benchmark results in CSV format.

### 4. Evaluation Methodology

Replicate paper benchmarks using identical datasets (mountain, street, terrain, building, campus scenes). Measure processing time, feature point counts, matching accuracy, and visual quality.

Compare LP-SIFT against OpenCV's SIFT/ORB/BRISK/SURF implementations + RANSAC.

### 5. Schedule of Work

**Phase 0** (Week 1: Nov 3-10): Run MATLAB baseline, validate paper results

**Phase 1** (Weeks 2-3: Nov 11-24): C++ implementation of LP-SIFT algorithm

**Owner:** Algorithms

- Image preprocessing, local-peak detection, SIFT descriptors, Feature matching, RANSAC homography, multi-image mosaic stitching
- **Deliverables:** `liblpsift.lib/lpsift_cli.exe` (pairwise & folder modes). Config file with parameters.

**Phase 2** (Week 4: Nov 25 - Dec 1): Benchmark replication

**Owner:** Evaluation

- Test on paper's datasets, compare with SIFT/ORB/BRISK/SURF
- **Deliverables:** `results.csv`, plots, side-by-side comparisons (C++ vs. MATLAB); table summarized by size bucket (small/medium/large).

**Phase 3** (Week 5: Dec 2-8): CPU optimizations

**Owner:** Optimization

- Implement C++ specific optimizations on the local peak algorithm (to be determined)
- **Deliverables:** Performance comparison with baseline C++ implementation

**Phase 3+** (Week 6: Dec 9-10): Final Report + GPU acceleration (stretch goal)

- CUDA kernels for gradient/histogram computation (The paper notes GPU/C++ would markedly improve efficiency over their MATLAB baseline.)
- **Deliverables:** Final report, demo video and/or live demonstration

### 6. Publicly Available Code and Data

**Reference Implementation:** MATLAB code from paper co-author Wei Zhao

**Baseline Comparisons:** OpenCV's cv::SIFT, cv::ORB, cv::BRISK, cv::SURF implementations

**Datasets:** Public benchmark images from paper references, custom 6MP smartphone captures  
**Our Contribution:** Optimized C++ port with SIMD/parallel optimizations and optional GPU acceleration, exceeding the original MATLAB implementation's performance

## 7. GitHub Repository

<https://github.com/davecyli/css587project>

Repository to include: C++ source code, library/CLI tool, test datasets, benchmark scripts/results, configuration files, documentation, and final report.

## 8. Preferred Dates

**Paper Presentation:** November 10 or 12, 2025

**Design Review:** November 24, 2025

## 9. Bibliography

- [1] Li, H., Wang, L., Zhao, T., & Zhao, W. (2024). Local-peak scale-invariant feature transform for fast and random image stitching. arXiv:2405.08578v2.
- [2] Lowe, D. G. (2004). Distinctive image features from scale-invariant keypoints. International Journal of Computer Vision, 60, 91-110.
- [3] Rublee, E., Rabaud, V., Konolige, K., & Bradski, G. (2011). ORB: An efficient alternative to SIFT or SURF. 2011 International Conference on Computer Vision, 2564-2571.
- [4] Bay, H., Ess, A., Tuytelaars, T., & Van Gool, L. (2008). Speeded-up robust features (SURF). Computer Vision and Image Understanding, 110(3), 346-359.
- [5] Leutenegger, S., Chli, M., & Siegwart, R. Y. (2011). BRISK: Binary robust invariant scalable keypoints. 2011 International Conference on Computer Vision, 2548-2555.
- [6] Fischler, M. A., & Bolles, R. C. (1981). Random sample consensus: A paradigm for model fitting. Communications of the ACM, 24(6), 381-395.